

Sharanya Krishnamurthi

AI Engineer

669-252-9725 | krishnamurthisharanya@gmail.com | Milpitas, CA, USA | linkedin.com/in/sharanya-krishnamurthi | sharanya-krishnamurthi.github.io/Portfolio/

EDUCATION

Georgia Institute of Technology <i>Masters in Computer Science, Artificial Intelligence Specialization</i>	Atlanta, GA Jan 2024 - May 2026
B. S. Abdur Rahman Crescent Institute of Technology <i>Bachelor of Technology in Electronics and Communication Engineering</i>	Chennai, India July 2015 - May 2019

EXPERIENCE

AI Chatbot Developer (Volunteer) — Health4TheWorld Inc <i>Developing AI and RAG-based Chatbots for Medical Use Cases</i>	Palo Alto, CA, USA Aug 2025 - Present
Member of Technical Staff - AI/ML	Zoho Corporation Pvt. Ltd., Chennai, India Oct 2019 - Nov 2023
<ul style="list-style-type: none">Designed and developed LLM-powered OCR pipelines for insurance EOB documents, achieving 95%+ field-level accuracy with reducing manual claims processing time by 70%.Built robust data preprocessing workflows, to address data drift and improve model accuracy by cleaning and de-identifying large-scale insurance datasets for downstream NLP and ML tasks.Trained LayoutLM model on the curated data using pytorch to achieve an average precision around 97% and a 94% average recall for all fields after 20 epochs on a NVIDIA A100 48GB GPUDeployed an ensemble model (rule-based + templates + LLMs) on Zoho's Stratus cloud and exposed it as a REST API, enabling seamless integration across products.Introduced a few-shot learning approach for line-item extraction using LayoutLM and Donut, achieving 60% automation while maintaining a 95%+ F1 score.Trained text detection and recognition models for diverse document types, enhancing OCR generalization.Acted as primary POC for document NLP initiatives, collaborating with product and engineering teams to deliver production-ready solutions.	

PROJECTS

AI-Powered Document QA System — RAG, Gemini 2.0, ChromaDB, Gradio, Hugging Face Spaces <ul style="list-style-type: none">Developed and deployed an end-to-end RAG system with Gemini 2.0, text-embedding-004, and ChromaDB, hosted on Hugging Face Spaces. Achieved 20% improvement in answer accuracy using prompt engineering.
Warehouse Search Optimization — Python, Robotics, Path Planning & Search, A* <ul style="list-style-type: none">Developed a warehouse automation solution to efficiently pick up and deliver boxes under different constraints.
Stock Portfolio Management using Machine Learning — Python, Decision Trees, BagLearner, Time-Series Analysis <ul style="list-style-type: none">Developed an ML-driven system to optimize stock portfolio performance using historical price data.
Agent for Solving Raven's Progressive Matrices — Python, Machine Learning, Computer Vision <ul style="list-style-type: none">Developed an intelligent agent capable of solving abstract reasoning problems similar to Raven's Progressive Matrices.Achieved 84/96 correct solutions (87.5%) , including 48 hidden problems, demonstrating strong pattern recognition.

TECHNICAL SKILLS

Languages: Python, Java, C/C++, SQL
Machine Learning & Deep Learning: PyTorch, Tensorflow, Keras, Scikit-learn, Hugging Face Transformers, Natural Language Processing (NLP), Computer Vision, Generative AI, Feature Engineering, Model Optimization
LLMs & Generative AI: Prompt Engineering, Fine-tuning, Retrieval-Augmented Generation (RAG), LangChain, Multi-modal Models, Large Language Models (LLMs)
Vector Databases & Knowledge Stores: ChromaDB, FAISS, Pinecone, AWS OpenSearch
MLOps & Cloud Platforms: AWS, GCP/Vertex AI, Docker, Kubernetes, Git, CI/CD Pipelines, Model Deployment, REST API Development, Experiment Tracking
Data Science Libraries: Pandas, NumPy, Matplotlib, Seaborn, OpenCV, SciPy

CERTIFICATIONS

AWS Certified AI Practitioner (2025)
